

SQL for Data Scientists

Data Science Retreat

Batch 44 - 2025

Krzysztof Buzar

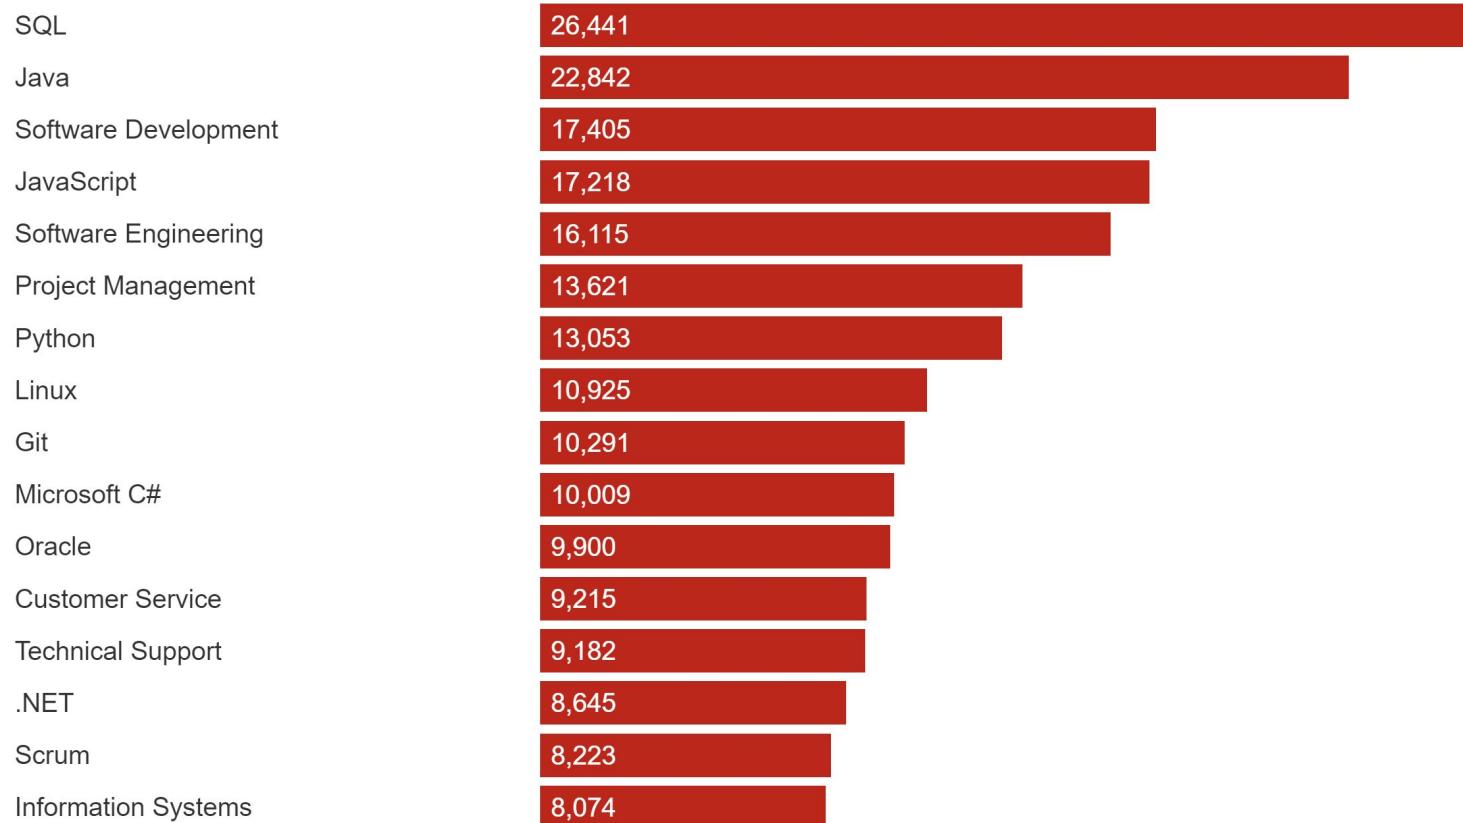
About me

- Python Developer
 - Background in academia (bioinformatics, computational physical chemistry)
 - Developed games (shortly)
 - Benchmarking enthusiast

Agenda

- Intro to SQL
 - What is this and why should we learn it?
 - Relational databases
 - Differences in SQL implementations
 - SQL vs NoSQL
- SQL syntax
 - SQLBolt drills
 - SQL review quiz
 - HackerRank drills
- Intro to Google BigQuery
- Python SQL interface
- Review questions

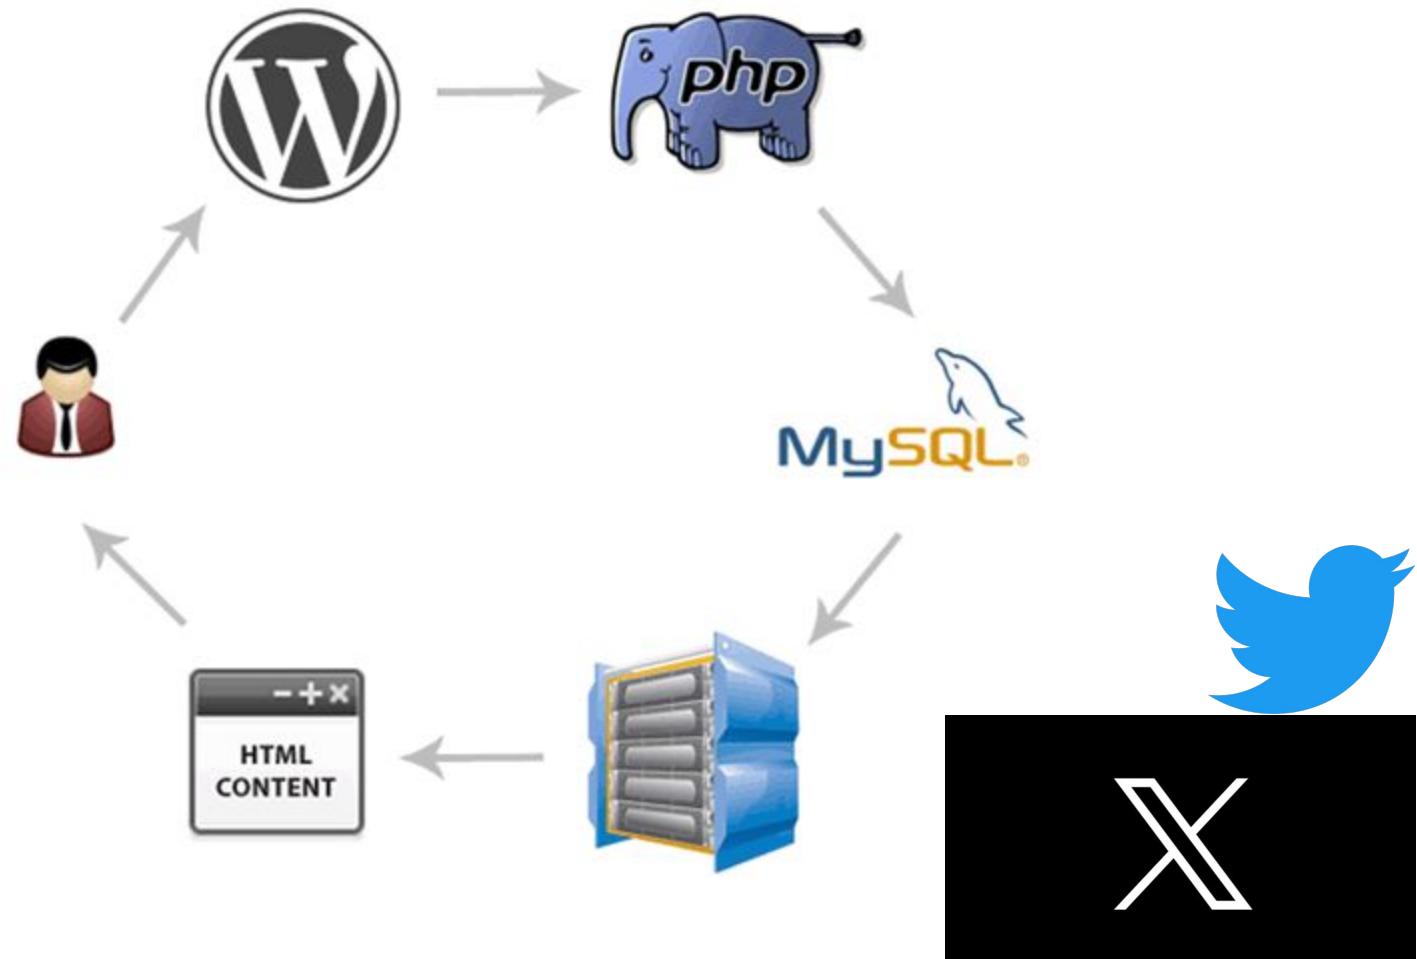
Most-Requested Tech Skills, Feb 18 - March 18, 2020



Source: Burning Glass Technologies (<https://www.burning-glass.com/>)

Hear of it before?

Who uses SQL?



SQL in different job roles

- ***Everyone*** in data uses SQL
- Data engineering, data analysis, machine learning engineering, software engineering, and data science roles all ***require some level of knowledge of SQL***

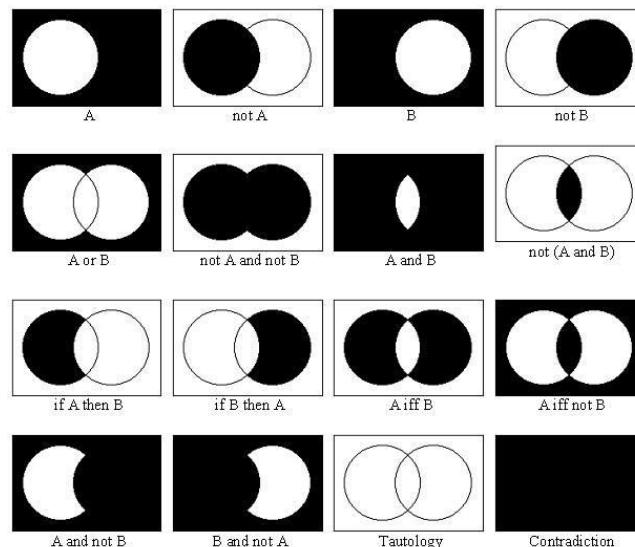
“

The popularity of SQL is not really surprising given the amount of data which is collected and churned out by thousands of businesses every single day. That said, it's only possible to make sense of big data if you have the right skillset. Data scientists rely heavily on SQL skills which is why SQL as a skill is in such high demand. Interestingly, in the 2019 edition of the Global Technical Hiring and Skills Report, SQL was the #1 IT most developers got tested regardless of their focus.

Source: Burning Glass Technologies (<https://www.burning-glass.com/>)

What is SQL?

- Pronounced “sequel” or “ES-CUE-L” (nobody should frown at either)
- Stands for “Structured Query Language” and is based in set theory
- Has been around since the 1970s (check out “[relational algebra](#)” in Wikipedia)



What is a relational database?

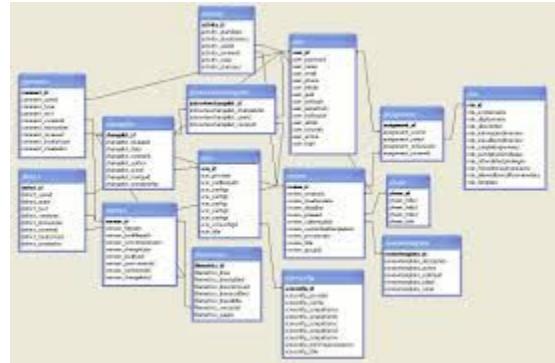


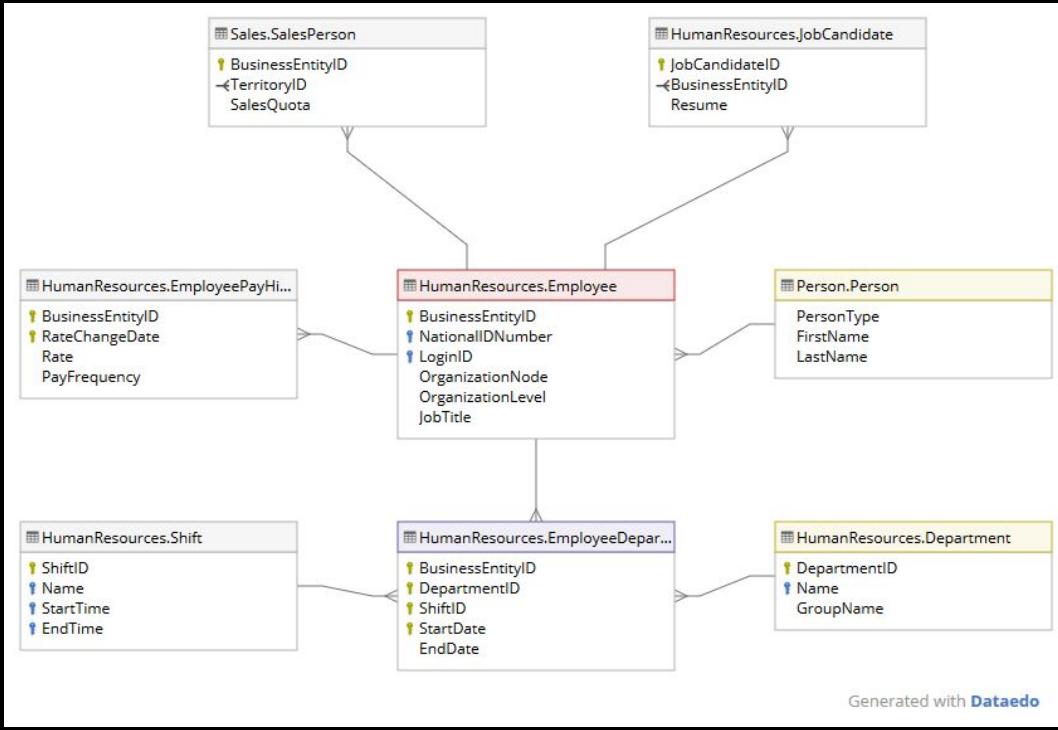
bodyweight + dumbbell workouts

	A	B	C	D	E
1		pushups	goblet squats	overhead press	squat + overhead press
2	27/04/20	198	85	20	0
3	28/04/20	128	36	0	0
4	29/04/20	72	0	0	0
5	30/04/20	276	0	0	0
6	01/05/20	0	100	0	0
7	02/05/20	0	0	0	0
8	03/05/20	67	25	0	0
9	04/05/20	96	25	0	0
10	05/05/20	108	75	0	0
11	06/05/20	0	0	0	0
12	07/05/20	66	0	0	0
13	08/05/20	622	111	0	0
14	09/05/20	15	0	0	0
15	10/05/20	0	0	0	0

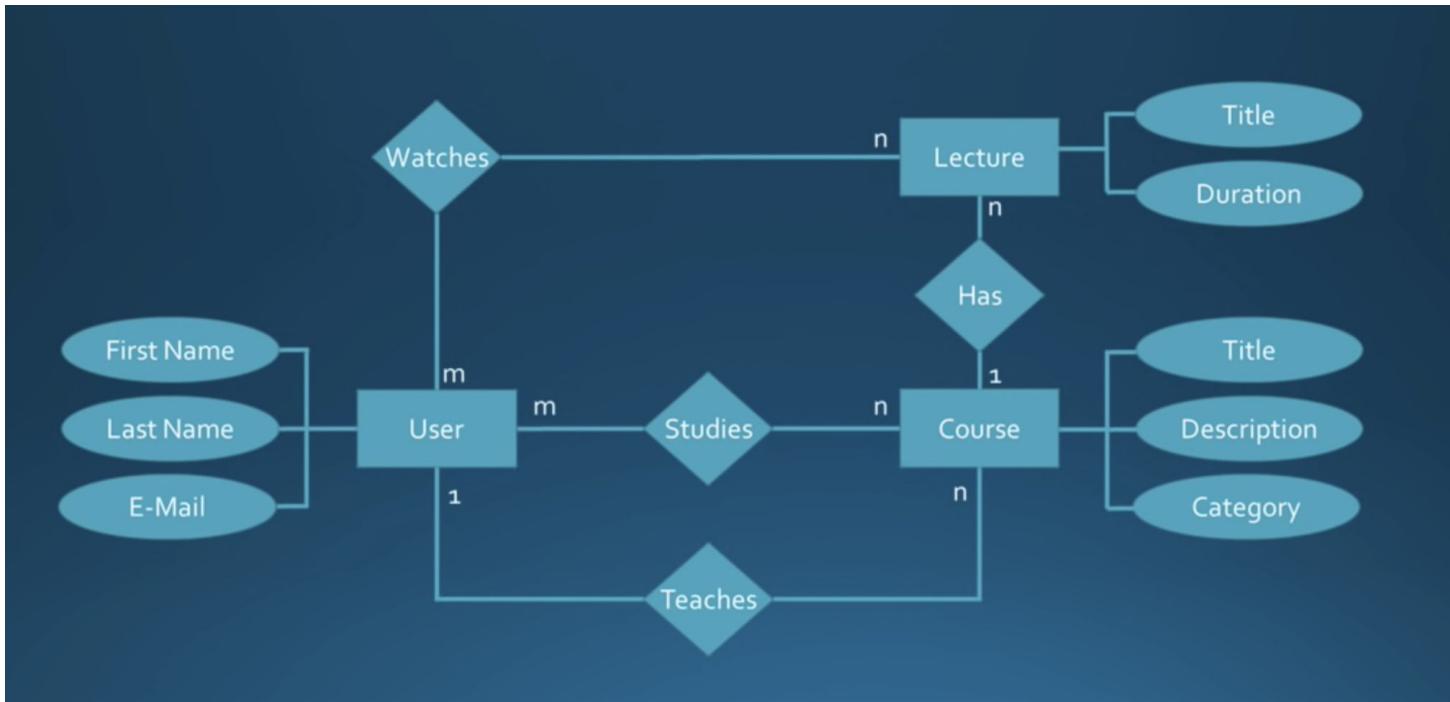
What is a schema?

A *schema* is a “*data model*”, the representation in tables of the **relationships** between **entities** (i.e., *nouns*, tables) and **attributes**(i.e. *adjectives*, features, columns within the tables)





Entity-Relationship diagrams

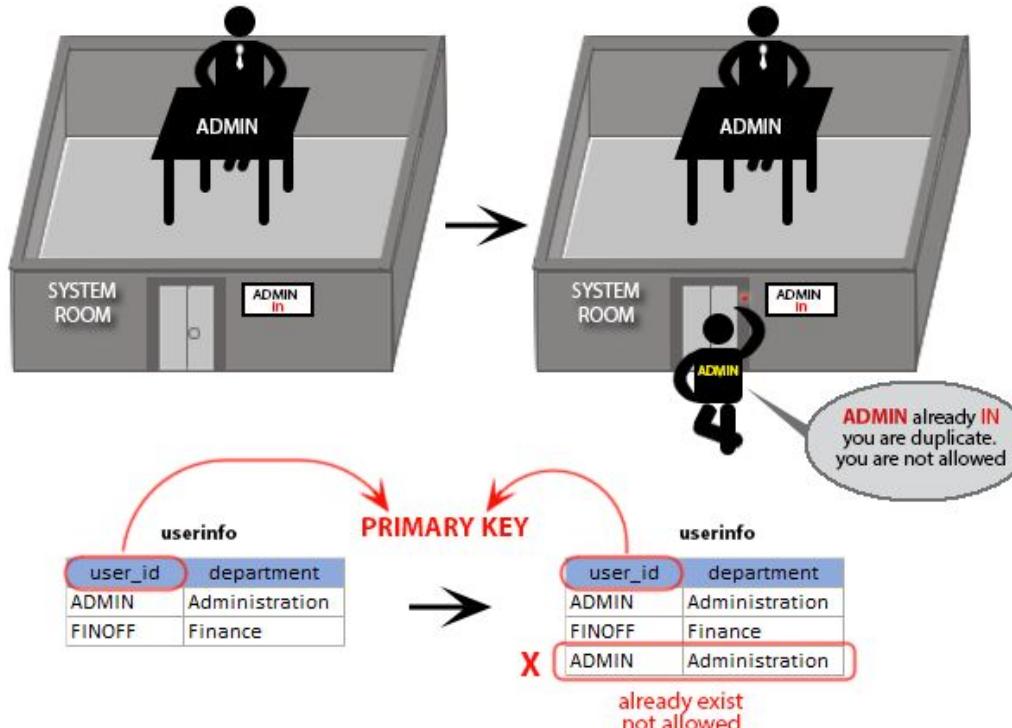


Source: [a brief tutorial on ER diagrams](#) (recommended view)

What's SQL role in an RDBMS?

- SQL was Created for managing data in relational database management systems (RDBMS)
- SQL takes *instances of relations* as input, yields them as output (views, tables)
- RDBMS comprise data models that detail relationships between tables
 - One to one
 - One to many
 - Many to many relationships

Understanding “primary keys”



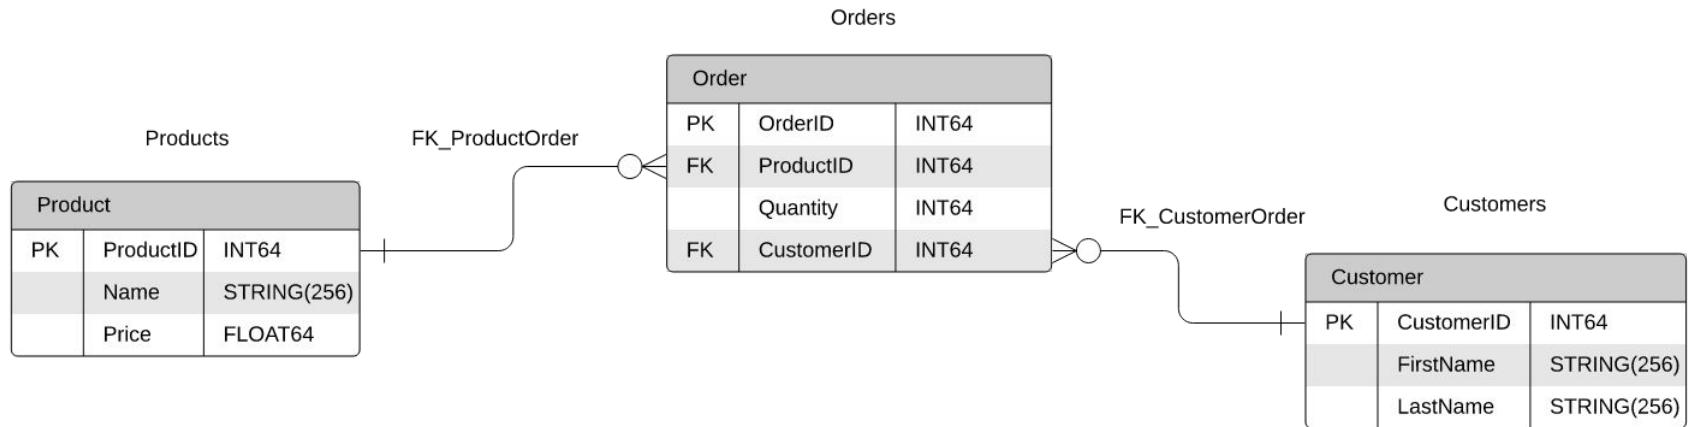
Understanding “primary keys”

They are called “keys” because they should be unique, as if “only one entity should have access to this resource (row)”

Commonly used primary keys:

- An automatically generated Universal Unique Identifier (UUID)
- An integer or hash value (for “internal” tables)
- A combination of two column values (read this on composite keys)
- ~~Social security / tax number.~~
- ~~An email address (but only if two users can't share the same email address)~~
- ~~Vehicle identification number.~~
- ~~Driver licence number~~
- Some other special code that is unique to each record

Understanding “foreign keys”



One to one relationships

Students table

Student ID	12345
Last Name	Tang
First Name	Sophie

Contact Info table

Student ID	12345
City	New York
Phone	408-555-3456

One to many relationships

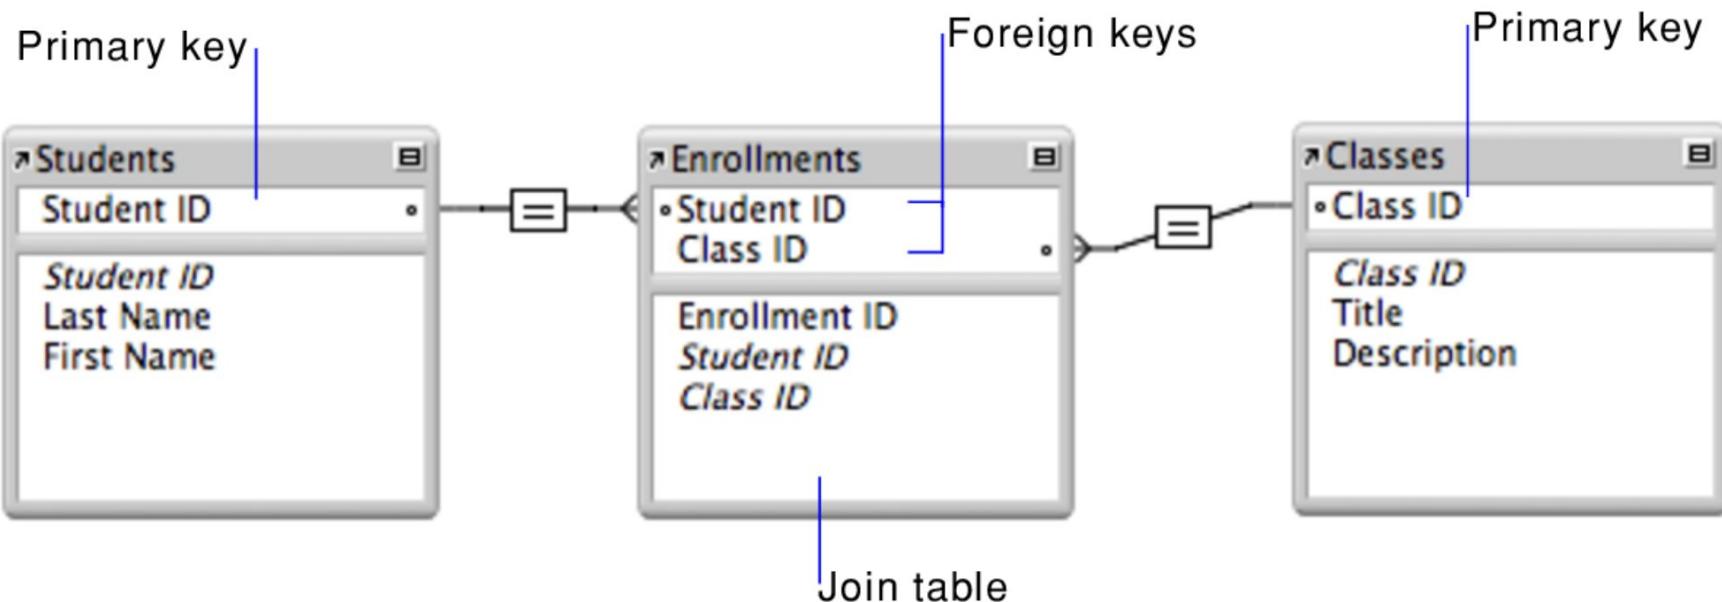
Customers table

Customer ID	12345
Name	Tang

Orders table

Order ID	B204
Customer ID	12345
Order ID	B391
Customer ID	12345
Order ID	B448
Customer ID	12345

Many to many relationships



Challenges of learning (and doing) SQL

- Learning SQL is boring
 - Possible solution: practice logic and syntax, write as many queries as possible
- Domain knowledge about a particular business is important to make reasonable data models
- We should **not** be learning SQL as if we were database administrators
 - Normalization and index optimization are outside of the scope of our jobs as data scientists
- Minor syntax differences between SQL implementations
 - Not a deal-breaker
- Understanding computational cost of queries is important while deciding on a SQL engine and while writing complex queries

The usefulness and beauty of SQLBolt



SQLBolt

Learn SQL with simple, interactive exercises.

- We won't need to install anything
- All answers are in the website
- Syntax is the same as in MySQL, Postgres, and SQLite
- 18 interactive beginner lessons, 2 intermediate / advanced ones on subqueries and set operators

<https://sqlbolt.com/>

Homework suggestion: Khan Academy's SQL

The image shows a screenshot of the Khan Academy website. At the top, there is a dark blue header bar with the following elements from left to right: a "Courses" dropdown menu, a search bar containing the word "Search" with a magnifying glass icon, and the Khan Academy logo, which consists of a green hexagon with a white stylized figure inside and the text "Khan Academy" next to it. Below the header, the main content area has a dark blue background. On the left side of this area, the text "Computer programming" is displayed in a light blue font. In the center, the words "Unit: Intro to SQL: Querying and managing data" are prominently displayed in large, bold, white font.

<https://www.khanacademy.org/computing/computer-programming/sql>

Optional: Install MySQL (Developer Settings)



- MySQL is the *most popular* relational database system
 - Used by Facebook, Twitter, Youtube, Flickr
 - “LAMP” stacks use it (it’s the ‘M’)
- Now owned by Oracle (many people don’t like this)
- MariaDB is a non-corporate open source fork

<https://dev.mysql.com/downloads/>



MySQL. Installer

Adding Community

Choosing a Setup Type

Installation

Installation Complete

Choosing a Setup Type

Please select the Setup Type that suits your use case.

Developer Default

Installs all products needed for MySQL development purposes.

Server only

Installs only the MySQL Server product.

Client only

Installs only the MySQL Client products, without a server.

Full

Installs all included MySQL products and features.

Custom

Manually select the products that should be installed on the system.

Setup Type Description

Installs the MySQL Server and the tools required for MySQL application development. This is useful if you intend to develop applications for an existing server.

This Setup Type includes:

* MySQL Server

* MySQL Shell

The new MySQL client application to manage MySQL Servers and InnoDB cluster instances.

* MySQL Router

High availability router daemon for InnoDB cluster setups to be installed on application nodes.

* MySQL Workbench

The GUI application to develop for and manage the server.

* MySQL for Excel

Next >

Cancel



MySQL® Installer

Adding Community

Choosing a Setup Type

Check Requirements

Installation

Product Configuration

Installation Complete

Check Requirements

The following products have failing requirements. MySQL Installer will attempt to resolve them automatically. Requirements marked as manual cannot be resolved automatically. Click on each item to try and resolve it manually.

For Product	Requirement	Status
<input checked="" type="radio"/> MySQL Server 8.0.20	Microsoft Visual C++ 2019 Redistributable Package (x64) is not installed. Latest binary compatible version will be installed if agreed	
<input type="radio"/> MySQL Workbench 8.0.20	Microsoft Visual C++ 2019 Redistributable Package (x64) is not installed. Latest binary compatible version will be installed if agreed	
<input type="radio"/> MySQL For Excel 1.3.8	Visual Studio 2010 Tools for Office Runtime is not installed. Latest binary compatible version will be installed if agreed	
<input type="radio"/> MySQL for Visual Studio 1.2.9	Visual Studio version 2015, 2017 or 2019 is required. Latest binary compatible version will be installed if agreed	Manual
<input type="radio"/> MySQL Shell 8.0.20	Microsoft Visual C++ 2019 Redistributable Package (x64) is not installed. Latest binary compatible version will be installed if agreed	
<input type="radio"/> MySQL Router 8.0.20	Microsoft Visual C++ 2019 Redistributable Package (x64) is not installed. Latest binary compatible version will be installed if agreed	

Requirement Details

MySQL Installer is trying to resolve this requirement automatically.
There is nothing you need to do.

Requirement:

Microsoft Visual C++ 2019 Redistributable Package (x64) is not installed. Latest binary compatible version will be installed if agreed

Status:

< Back

Execute

Next >

Cancel

 MySQL®. Installer

Adding Community

Choosing a Setup Type

Installation

Product Configuration

Installation Complete

Installation

The following products will be installed.

Product	Status	Progress	Notes
 MySQL Notifier 1.1.8	Complete		
 Connector/ODBC 8.0.20	Failed		
 Connector/C++ 8.0.20	Complete		
 Connector/J 8.0.20	Complete		
 Connector/.NET 8.0.20	Complete		
 Connector/Python 8.0.20	Failed		
 MySQL Documentation 8.0.20	Complete		
 Samples and Examples 8.0.20	Complete		

Show Details >

< Back

Next >

Cancel

MySQL Installer

X



MySQL Installer

No compatible servers were found. You'll need to cancel this wizard and install one.

OK

Python and MySQL setup

Recommended tutorial:

<https://realpython.com/python-mysql/>

What is normalization?

- Think of it as “**database cleanup**”
- Normalization **reduces duplicated/redundant/mismatched data**
- Usually a job for **database administrators**, not data scientists
- Curious? https://en.wikipedia.org/wiki/Database_normalization

Normalization as redundancy minimization

- Repetition of data increases the size of the database
 - Minor problem
- Data integrity issues
 - Insertion
 - Deletion
 - Update
- Suppose we have a table shared info and repeated info among columns
 - The repeated info can be put into a separate table
 - To query the database we need to JOIN with the table that avoids duplication

Normalization as redundancy minimization

STUDENTS TABLE				
rollno	name	branch	hod	office_tel
1	Akon	CSE	Mr. X	53337
2	Bkon	CSE	Mr. X	53337
3	Ckon	CSE	Mr. X	53337
4	Dkon	CSE	Mr. X	53337

Normalization as redundancy minimization

STUDENTS TABLE				
rollno	name	branch	hod	office_tel
1	Akon	CSE	Mr. X	53337
2	Bkon	CSE	Mr. X	53337
3	Ckon	CSE	Mr. X	53337
4	Dkon	CSE	Mr. X	53337

Data integrity issues and computational cost while updating is to be considered

<https://www.youtube.com/watch?v=xoTyrdT9SZII>

Normalization as redundancy minimization

STUDENTS TABLE			BRANCH TABLE		
rollno	name	branch	branch	hod	office_tel
1	Akon	CSE	CSE	Mr. Y	53337
2	Bkon	CSE			
3	Ckon	CSE			

Splitting tables allows us to reduce the number of required updates among rows

<https://www.youtube.com/watch?v=xoTyrdT9S1>

Normalization as redundancy minimization

STUDENTS TABLE			BRANCH TABLE		
rollno	name	branch	branch	hod	office_tel
			CSE	Mr. Y	53337

No Data



Q: What would happen if we were to put all of our data into a single table?

<https://www.youtube.com/watch?v=xoTyrdT9SZ>

First Normal Form (1NF)

It is **poor design** if a table is not in 1NF

1. Columns contain atomic values, not lists of values
2. Columns contain values of the same type
3. Columns should have unique names
4. Row order should not matter

STUDENTS TABLE		
rollno	name	subject
101	Akon	OS, CN
103	Ckon	JAVA
102	Bkon	C, C++

STUDENTS TABLE		
rollno	name	subject
101	Akon	OS
101	Akon	CN
103	Ckon	JAVA
102	Bkon	C
102	Bkon	C++

2NF, 3NF, BCNF, 4NF, 5NF

2ND and 3NF are applied on partial and transitive dependencies on composite keys

- [2NF](#)
- [3NF](#)
- [BCNF](#)
- [4NF](#)
- [5NF](#)

Using UUIDS as primary keys eliminates the need for most of these rules

- **This is a job for database administrators, not data scientists**

On different SQL implementations

‘Standard SQL’ syntax has about 90% overlap between different implementations:

- MySQL/MariaDB
- PostgreSQL
- Oracle
- Microsoft’s SQL Server
- IBM’s Dbase

SQL syntax differences between implementations

Operators	Works in
Arithmetic operators	All databases
Assignment operators	All databases
Bitwise operators	Microsoft SQL Server
Comparison operators	All databases
Logical operators	DB2, Oracle, SQL Server, and PostgreSQL
Unary operators	DB2, Oracle, and SQL Server

SQL vs NoSQL

- NoSQL databases are (e.g. MongoDB, DynamoDB, Cassandra) are **write-optimized**
 - Easily mutable
 - Reduced need to update/maintain a schema
- SQL is **read-optimized**, good for querying and understanding data relationships
 - Easily queryable **for a data analyst**
 - Needs maintenance of a schema

[A nice video on SQL vs NoSQL](#)

** End users won't be able to tell the difference between these systems. Only backend engineers and data people care about this*

SQL vs NoSQL

SQL: uses ‘structured’ data (aka tables)

Strengths of SQL

- Read optimized
- Easy querying
- Think of this as standard analysis on a curated dataset of clinical data

NoSQL: uses ‘non-structured’ data (aka not tables, things like JSON)

Strengths of NoSQL

- Write optimized
- Parallel concurrent writing
- Think of this as what Spotify needs to do when they capture your usage data

RELATIONAL DATABASE



DATA LAKE



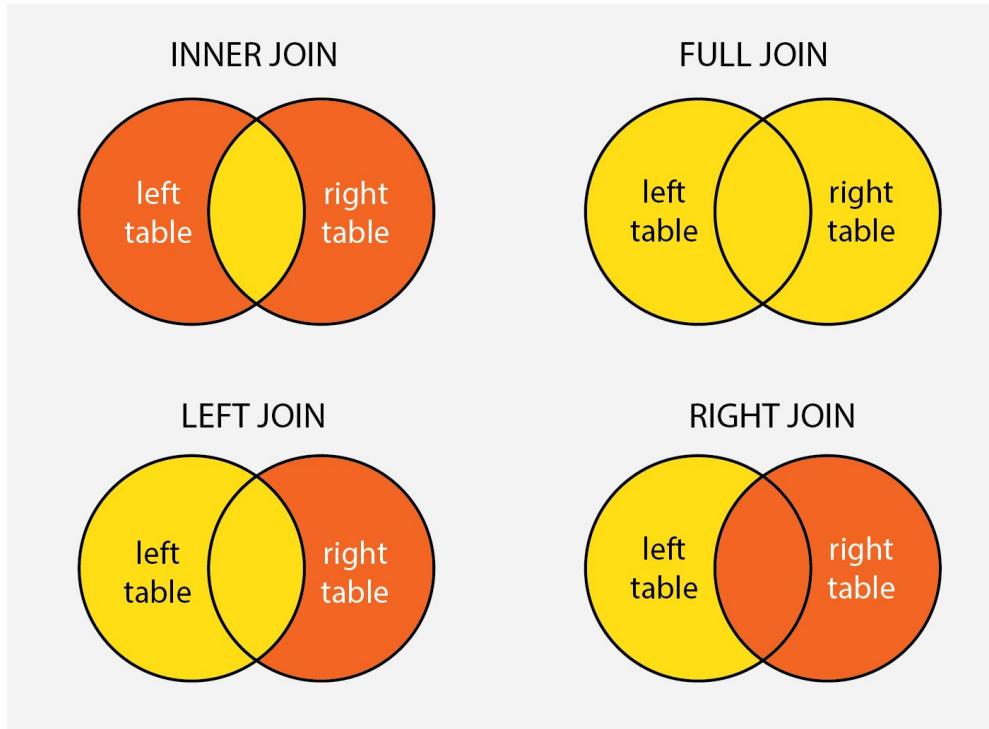
W3Schools: a doubleplusgood SQL syntax guide

w3schools.com

SQL Tutorial
SQL HOME
SQL Intro
SQL Syntax
SQL Select
SQL Select Distinct
SQL Where
SQL And, Or, Not
SQL Order By
SQL Insert Into
SQL Null Values
SQL Update
SQL Delete
SQL Select Top
SQL Min and Max
SQL Count, Avg, Sum
SQL Like

<https://www.w3schools.com/sql/>

Types of Joins

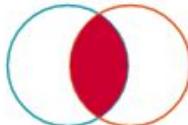


Types of Joins

Visual JOIN

Understand how joins work by interacting and see it visually

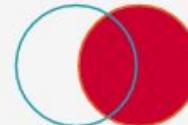
INNER JOIN (or JOIN)



LEFT JOIN



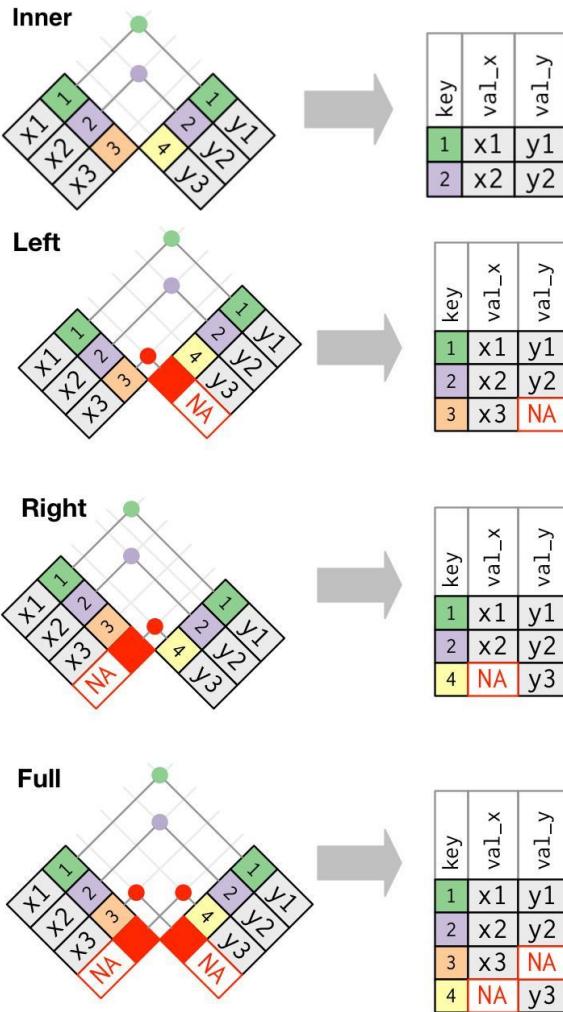
RIGHT JOIN

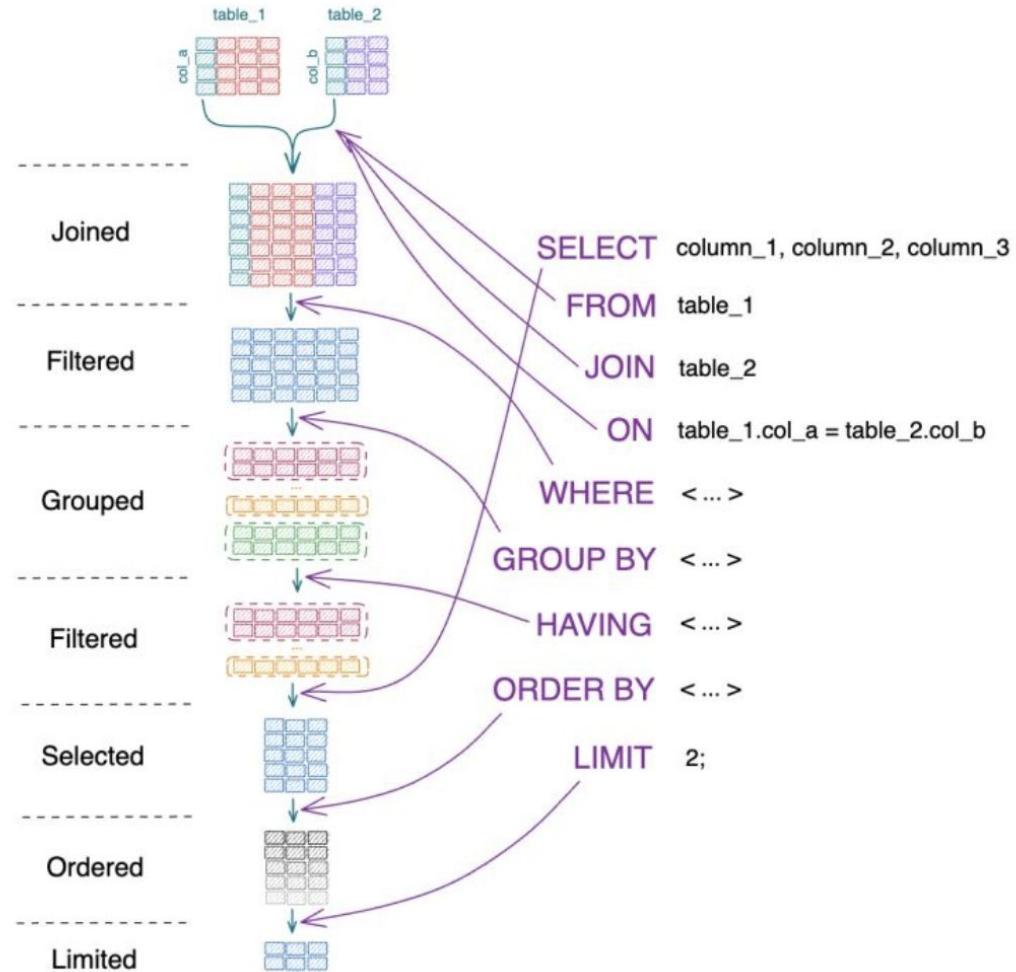


OUTER JOIN (with UNION)



[Visual JOIN](#)





Understanding the inner vs outer joins

martian_id	first_name	last_name	base_id	super_id
13	John	Carter	null	8



base_id	base_name	founded
1	Tharsisland	2037-06-03
2	Valles Marineris 2.0	2040-12-01
3	Gale Cratertown	2041-08-15
4	New New New York	2042-02-10
5	Olympus Mons Spa & Casino	null

martian_id	first_name	last_name	base_id	super_id	base_id	base_name	founded
1	Ray	Bradbury	1	null	1	Tharsisland	2037-06-03
2	John	Black	4	10	4	New New New York	2042-02-10
3	Samuel	Hinkston	4	2	4	New New New York	2042-02-10
4	Jeff	Spender	1	9	1	Tharsisland	2037-06-03
5	Sam	Parkhill	2	12	2	Valles Marineris 2.0	2040-12-01
6	Elma	Parkhill	3	8	3	Gale Cratertown	2041-08-15
7	Melissa	Lewis	1	1	1	Tharsisland	2037-06-03
8	Mark	Watney	3	null	3	Gale Cratertown	2041-08-15
9	Beth	Johanssen	1	1	1	Tharsisland	2037-06-03
10	Chris	Beck	4	null	4	New New New York	2042-02-10
11	Nathaniel	York	4	2	4	New New New York	2042-02-10
12	Elon	Musk	2	null	2	Valles Marineris 2.0	2040-12-01

Going through SQLBolt

- **Rules/order of this honor game**
 1. We read the theory page of the lesson together
 2. Each one of us attempts to write each query individually. If it takes more than 2 minutes, we ask a peer
 3. If two of us are amiss, we look at the solution
 4. We send a Slack message saying that we are done with the lesson
 5. Ask as many questions as you like, but first *try doing the exercise*

How does a SQL query look like? SQLBolt lesson 1

Table: Movies

<u>Id</u>	<u>Title</u>	<u>Director</u>	<u>Year</u>	<u>Length_minutes</u>
1	Toy Story	John Lasseter	1995	81
2	A Bug's Life	John Lasseter	1998	95
3	Toy Story 2	John Lasseter	1999	93
4	Monsters, Inc.	Pete Docter	2001	92
5	Finding Nemo	Andrew Stanton	2003	107
6	The Incredibles	Brad Bird	2004	116
7	Cars	John Lasseter	2006	117
8	Ratatouille	Brad Bird	2007	115
9	WALL-E	Andrew Stanton	2008	104
10	Up	Pete Docter	2009	101

The query to get everything in a table

```
SELECT *
FROM mytable;
```

This query retrieves everything - no restrictions

Selecting columns - SQLBolt lesson 1

```
SELECT column, another_column  
FROM mytable;
```

A query is a selection of data. In SQL, anything that starts with a SELECT statement is a query.

Table: Movies

Title	Director
Toy Story	John Lasseter
A Bug's Life	John Lasseter
Toy Story 2	John Lasseter
Monsters, Inc.	Pete Docter
Finding Nemo	Andrew Stanton
The Incredibles	Brad Bird
Cars	John Lasseter
Ratatouille	Brad Bird
WALL-E	Andrew Stanton
Up	Pete Docter

```
SELECT title,director FROM movies;
```

Truth tables for AND, OR, and NOT

A	B	A AND B	A OR B	NOT A
False	False	False	False	True
False	True	False	True	True
True	False	False	True	False
True	True	True	True	False

Queries with constraints

Operator	Condition	SQL Example
=, !=, < <=, >, >=	Standard numerical operators	col_name != 4
BETWEEN ... AND ...	Number is within range of two values (inclusive)	col_name BETWEEN 1.5 AND 10.5
NOT BETWEEN ... AND ...	Number is not within range of two values (inclusive)	col_name NOT BETWEEN 1 AND 10
IN (...)	Number exists in a list	col_name IN (2, 4, 6)
NOT IN (...)	Number does not exist in a list	col_name NOT IN (1, 3, 5)

Working with aggregates

Function	Description
COUNT(*) , COUNT(column)	A common function used to counts the number of rows in the group if no column name is specified. Otherwise, count the number of rows in the group with non-NULL values in the specified column.
MIN(column)	Finds the smallest numerical value in the specified column for all rows in the group.
MAX(column)	Finds the largest numerical value in the specified column for all rows in the group.
AVG(column)	Finds the average numerical value in the specified column for all rows in the group.
SUM(column)	Finds the sum of all numerical values in the specified column for the rows in the group.

Working with aggregates

Role	Years_employed
Artist	7
Engineer	1
Manager	6

```
SELECT Role, Years_employed FROM employees GROUP BY Role;
```

Exercise 10 — Tasks

1. Find the longest time that an employee has been at the studio ✓
2. For each role, find the average number of years employed by employees in that role
3. Find the total number of employee years worked in each building

Stuck? Read this task's [Solution](#).

Solve all tasks to continue to the next lesson.

What's missing in this query?

Working with aggregates

Building	Years_employed
1e	3
2w	6

```
SELECT Building, Years_employed| FROM Employees GROUP BY  
Building
```

Exercise 10 — Tasks

1. Find the longest time that an employee has been at the studio ✓
2. For each role, find the average number of years employed by employees in that role
3. Find the total number of employee years worked in each building

Stuck? Read this task's [Solution](#).
Solve all tasks to continue to the next lesson.

What's missing in this query?

Using DISTINCT on aggregate functions

```
SELECT COUNT(DISTINCT Customer_Name)  
FROM Customers
```

If **DISTINCT is not specified, ALL is assumed*

HAVING as a constraint for aggregates

```
SELECT group_by_column, AGG_FUNC(column_expression) AS aggregate  
FROM mytable  
WHERE condition  
GROUP BY column  
HAVING group_condition;
```

After filtering with WHERE, we can apply a second filter using HAVING

Leetcode - left join example

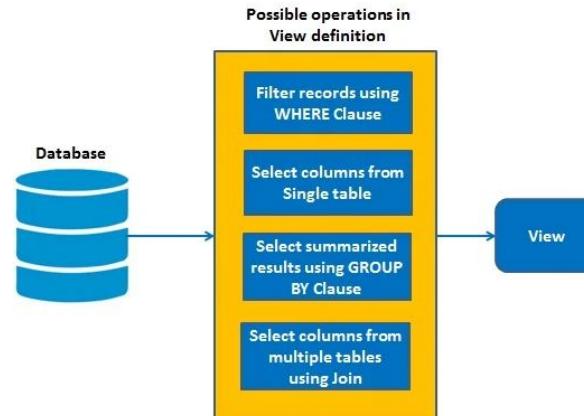
<https://leetcode.com/problems/combine-two-tables/>

Solution (Highlight text to see) :

Queries with views

Views are tables that don't persist in a hard drive (views are stored as commands)

- They simplify complex queries
- Prevent db users from accessing sensitive columns



Homework for practicing SQL Joins: W3Schools

Test Yourself With Exercises

Exercise:

Insert the missing parts in the `JOIN` clause to join the two tables `Orders` and `Customers`, using the `CustomerID` field in both tables as the relationship between the two tables.

```
SELECT *
FROM Orders
LEFT JOIN Customers
    [ ] = [ ];
;
```

[Submit Answer >](#)

https://www.w3schools.com/sql/sql_join.asp

(Basic) SQL Review Quiz

Question 1 of 25:

What does SQL stand for?

- Structured Question Language
- Structured Query Language
- Strong Question Language

Next ›

0

<https://www.w3schools.com/quiztest/quiztest.asp?qtest=SQL>

Style suggestions: making your SQL readable

- Names must begin with a letter and may not end with an underscore
- Never give a table the same name as one of its columns
- Avoid plurals—use the more natural collective term where possible instead. For example `staff` instead of `employees` or `people` instead of `individuals`.
- Avoid abbreviations and if you have to use them make sure they are commonly understood

<https://www.sqlstyle.guide/>

Subqueries

Subqueries are *queries within queries (can easily become unreadable)*

```
SELECT *
FROM sales_associates
WHERE salary >
    (SELECT AVG(revenue_generated)
     FROM sales_associates);
```

SQL Syntax - Exercises on HackerRank

Higher Than 75 Marks

Easy, Max Score: 15, Success Rate: 98.78%



Solved

Employee Names

Easy, Max Score: 10, Success Rate: 99.72%



Solved

Employee Salaries

Easy, Max Score: 10, Success Rate: 99.64%



Solved

Average Population of Each Continent

Easy, Max Score: 10, Success Rate: 97.75%



Solved

The Report

Medium, Max Score: 20, Success Rate: 96.80%



Solved

[Solutions](#)

What is Google Bigquery?

- Uses standard SQL syntax
- Runs queries on petabytes of data in seconds
- No need to create or manage indices to speed up queries
- No need to maintain infrastructure

BigQuery sandbox

The screenshot shows the Google Cloud Platform BigQuery interface. On the left, there's a sidebar with links for Query history, Saved queries, Job history, Transfers, Scheduled queries, Reservations, BI Engine, and Resources. A search bar at the top right contains the text "test-bigquery". The main area is titled "Unsaved query" and shows the following SQL query:

```
1 SELECT author, text FROM `bigquery-public-data.hacker_news.comments` LIMIT 1000
```

Below the query are buttons for Run, Save query, Save view, Schedule query, More, Hide Editor, and Full Screen. A note indicates that the query will process 3 GB when run. The table being queried is named "comments". The "Preview" tab is selected, showing the schema (Row, id, by, author, time, time_ts, text) and three rows of data. The first row is: Row 1, id 2701393, by 5l, author 5l, time 1309184881, time_ts 2011-06-27 14:28:01 UTC, text "And the glazier who fixed all the broken windows also left his money to good causes.". The second row is: Row 2, id 5811403, by 99, author 99, time 1370234048, time_ts 2013-06-03 04:34:08 UTC, text "Does canada have the equivalent of H1B/Green card for work sponsorship? What do you". The third row is: Row 3, id 21623, by AF, author AF, time 1178992400, time_ts 2007-05-12 17:53:20 UTC, text "Speaking of Rails, there are other options in the Python world besides Django.<p>Pylons i". A warning message states that some cell values are very long and are truncated to the first 1024 characters. A "Show full values" button is available. At the bottom, there are buttons for Rows per page (set to 100), First page, Last page, and Next/Previous page.

Using the BigQuery sandbox

[Query without a credit card: introducing BigQuery sandbox](#)

Suggested homework: Kaggle's Bigquery Tutorial

Intro to SQL

Learn SQL for working with databases, using Google BigQuery.

View Certificate



100% complete! Congrats!

<https://www.kaggle.com/learn/intro-to-sql#>

Suggested homework: Google's BigQuery



[Back to CIFL](#)

[Dashboard](#)

AR ▾

Getting Started with BigQuery

Making the jump from spreadsheets to databases,
one small SQL query at a time.

[Resume course](#)

```
1 SELECT
2 date,
3 crawl_date,
4 url,
5 top_admin_action,
6 top_admin_action_reason
7 FROM `data-pipeline-app.wga_codingisforlosers.actions_data_studio`
8 where date = cast('2019-07-01' as date)
9 and top_admin_action != ''
10 and url = 'codingisforlosers.com/google-sheets-query-function'
11 order by sessions_30d desc
```

Standard SQL Dialect

[RUN QUERY](#) ▾ [Save Query](#) [Save View](#) [Format Query](#) [Schedule Query](#) [Show Options](#)

Results

Results				
Row	date	crawl_date	url	top_admin_action
1	2019-07-01	2019-08-30	codingisforlosers.com/google-sheets-query-function	missing canonical

<https://learn.codingisforlosers.com/learn-bigquery-sql>

References and further reading

- SQL interview questions
 - <https://github.com/alexeygrigorev/data-science-interviews/blob/master/technical.md#sql>
- SQL for Data Science - Coursera
 - <https://www.coursera.org/learn/sql-for-data-science>
- Kaggle's Intro to SQL with Bigquery
 - <https://www.kaggle.com/learn/intro-to-sql>
- Khan Academy
 - <https://www.khanacademy.org/computing/computer-programming/sql>
- BigQuery Tutorial - CIFL
 - <https://learn.codingisforlosers.com/learn-bigquery-sql>
- Understanding database normalization
 - <https://medium.com/cracking-the-data-science-interview/an-introduction-to-big-data-data-normalization-b72311f134b7>
- Common mistakes when working with dbs
 - <https://www.craigkerstiens.com/2016/06/07/five-mistakes-databases/>

Review questions

- What is a data model? What is a schema?
- Why are databases divided into tables?
- What is the difference between a primary key and a foreign key?
- What are the considerations of using a German social security number as primary key a table called Employee?
- After using GROUP BY, do we filter results with WHERE?
- SQL vs NoSQL: if a database requires real-time updating of its data model (e.g. adding more columns and/or changing the relationships between them) do we say that it should be ‘write-optimized’ or ‘read-optimized’?

Review questions

- How would you explain normalization? What is 1NF?
- If a SQL database query is longer (i.e. has more code instructions written in SQL), does that mean that it will take longer to run?
- What does an inner join do? What does a left join do?
- What considerations should we have if we store a jpg image as a row in a database? (Consider the BLOB type in MySQL) When would you prefer storing this resource elsewhere and keeping an URL?